

Oat is a set of programs for processing images, extracting object position information, and streaming data to disk and/or the network in real-time. Oat subcommands are independent programs that each perform a single operation but that can communicate through shared memory. This allows a user to chain operations together in arrangements suitable for particular context or tracking requirement. This architecture enables scripted construction of custom data processing chains. Oat is primarily used for real-time animal position tracking in the context of experimental neuroscience, but can be used in any circumstance that requires real-time object tracking.



Contributors

- jonnew <http://www.mit.edu/~jpnewman/>

Table of Contents

- Manual
 - Introduction
 - Frame Server
 - * Signature
 - * Usage
 - * Configuration File Options
 - * Examples
 - Frame Filter
 - * Signature
 - * Usage
 - * Configuration File Options
 - * Examples
 - Frame Viewer
 - * Signature
 - * Usage
 - * Example
 - Position Detector
 - * Signature
 - * Usage
 - * Configuration File Options
 - * Example
 - Position Generator
 - * Signature
 - * Usage

- * Configuration File Options
 - * Example
- Position Filter
 - * Signature
 - * Usage
 - * Configuration File Options
 - * Example
- Position Combiner
 - * Signature
 - * Usage
 - * Configuration File Options
 - * Example
- Frame Decorator
 - * Signature
 - * Usage
 - * Example
- Recorder
 - * Signature
 - * Usage
 - * Example
- Position Network Socket
 - * Signature
 - * Usage
 - * Example
- Buffer
 - * Signatures
 - * Usage
 - * Example
- Calibrate
 - * Signature
 - * Usage
- Kill
 - * Usage
 - * Example
- Clean
 - * Usage
 - * Example
- Installation
 - Dependencies
 - * Flycapture SDK

- * Boost
- * OpenCV
- * RapidJSON, cpptoml, and Catch
- * Setting up a Point-grey PGE camera in Linux
 - Camera IP Address Configuration
 - Point Greg GigE Host Adapter Card Configuration
 - Multiple Cameras
 - Example
- TODO

Manual

Introduction

Oat's design is influenced by the UNIX philosophy, suckless tools, and MEABench. Oat consists of a set of small, composable programs (called **components**). Components are equipped with standard interfaces that permit communication through shared memory to capture, process, and record video streams. Currently, Oat components act on two basic data types: **frames** and **positions**.

- **frame** - Video frame.
- **position** - 2D position.

Oat components can be chained together to realize custom dataflow networks that operate on instances of the aforementioned datatypes, called **tokens**. Token processing pipelines can be split and merged while maintaining thread-safety and sample synchronization. The messaging library underlying the communication between Oat components has been optimized to reduce token copying. For instance, **frame** passing is performed using a zero-copy protocol. This means that passing **frames** between components in a user-configured processing network incurs almost no memory or CPU cost compared to the monolithic equivalent. Further, great care was taken during implementations of Oat components to minimize time spent in critical sections. This means that individual components execute largely in parallel, even when components are highly interdependent, facilitating efficient use of multi-core CPUs and GPU-based processing acceleration.

To get a feel for how Oat is used, here is a script to detect the position of a single object in pre-recorded video file:

```
# Serve frames from a video file to the 'raw' stream
oat frameserve file raw -f ./video.mpg &

# Perform background subtraction on the 'raw' stream
# Serve the result to the 'filt' stream
# If an appropriately configured GPU is available, this process will
# use it
oat framefilt mog raw filt &

# Perform color-based object position detection on the 'filt' stream
# Serve the object position to the 'pos' stream. Allow parameter tuning
# through a simple GUI.
oat posidet hsv filt pos --tune &

# Decorate the 'raw' stream with the detected position from the 'pos' stream
# Serve the decorated images to the 'dec' stream
```

```

oat decorate -p pos raw dec &

# View the 'dec' stream
oat view dec &

# Record the 'dec' and 'pos' streams to file in the current directory
oat record -i dec -p pos -f ./

```

This script has the following graphical representation:

```

frameserve --> framefilt --> posidet --> decorate ---> view
      \                               /      \
      -----> record

```

Generally, an Oat component is called in the following pattern:

```
oat <subcommand> [TYPE] [IO] [CONFIGURATION]
```

The **<subcommand>** indicates the component that will be executed. Components are classified according to their type signature. For instance, **framefilt** (frame filter) accepts a frame and produces a frame. **posifilt** (position filter) accepts a position and produces a position. **frameserve** (frame server) produces a frame, and so on. The **TYPE** parameter specifies a concrete type of transform (e.g. for the **framefilt** subcommand, this could be **bsub** for background subtraction). The **IO** specification indicates where the component will receive data from and to where the processed data should be published. The **CONFIGURATION** specification is used to provide parameters to shape the component's operation. Aside from command line options and switches, which are listed using the **--help** option for each subcommand, the user can often provide an external file containing a configuration table to pass parameters to a component. Some configuration parameters can only be specified using a configuration file. Configuration files are written in plain text using TOML. A multi-component processing script can share a configuration file because each component accesses parameter information using a file/key pair, like so

```

[key]
parameter_0 = 1           # Integer
parameter_1 = true        # Boolean
parameter_2 = 3.14         # Double
parameter_3 = [1.0, 2.0, 3.0] # Array of doubles

```

or more concretely,

```
# Example configuration file for frameserve --> framefilt
[frameserve-config]
frame_rate = 30                # FPS
roi = [10, 10, 100, 100]      # Region of interest

[framefilt-config]
mask = "~/Desktop/mask.png"    # Path to mask file
```

The type and sanity of parameter values are checked by Oat before they are used. Below, the type signature, usage information, available configuration parameters, examples, and configuration options are provided for each Oat component.

Frame Server

`oat-frameserve` - Serves video streams to shared memory from physical devices (e.g. webcam or GIGE camera) or from file.

Signature

`oat-frameserve --> frame`

Usage

Usage: `frameserve` [INFO]
or: `frameserve` TYPE SINK [CONFIGURATION]
Serve image stream to a frame SINK

TYPE:

`wcam`: Onboard or USB webcam.
`gige`: Point Grey GigE camera.
`file`: Video from file (*.mpg, *.avi, etc.).
`test`: Write-free static image server for performance testing.

SINK:

User-supplied name of the memory segment to publish frames to (e.g. `raw`).

OPTIONAL ARGUMENTS:

INFO:

`--help` Produce help message.
`-v [--version]` Print version information.

CONFIGURATION:

`-i [--index] arg` Index of camera to capture images from.
`-f [--file] arg` Path to video file if 'file' is selected as the server TYPE.
 Path to image file if 'test' is selected as the server TYPE.
`-r [--fps] arg` Frames per second. Overriden by information in configuration file if provided.
`-c [--config] arg` Configuration file/key pair.

Configuration File Options

TYPE = `gige`

- **index=+int** User specified camera index. Useful in multi-camera imaging configurations.
- **fps=+float** Acquisition frame rate (Hz). Ignored if **trigger_on=true** and **enforce_fps=false**. If unspecified, then the maximum frame rate will be used.
- **exposure=float** Automatically adjust both shutter and gain to achieve given exposure (EV).
- **shutter=+float** Shutter time in milliseconds. Specifying **exposure** overrides this option.
- **gain=float** Sensor gain value. Specifying **exposure** overrides this option (dB).
- **white_bal={red=+int, blue=+int}** White-balance specified as red/blue intensity values (0-1000).
- **roi={x_offset=+int, y_offset=+int, width=+int, height=+int}** Region of interest to extract from the camera or video stream (pixels).
- **trigger_on=bool** True to use camera trigger, false to use software polling.
- **trigger_rising=bool** True to trigger on rising edge, false to trigger on falling edge.
- **trigger_mode=+int** Point-grey trigger mode. Common values are:
 - 0 - Standard external trigger. Trigger edge causes sensor exposure, then sensor readout to internal memory.
 - 1 - Blub shutter mode. Same as 0, except that sensor exposure duration is determined by trigger active duration.
 - 7 - Continuous internal trigger. No external trigger required, but not synchronized to an external clock.
 - 14 - Overlapped exposure/readout external trigger. Sensor exposure occurs during sensory readout to internal memory. This is the fastest external trigger mode.
- **trigger_pin=+int** Hardware pin number on Point-grey camera that trigger is sent to.
- **strobe_pin=+int** Hardware pin number on Point-grey camera that a gate signal for the camera shutter is copied to.
- **enforce_fps=bool** If true, ensures that frames are produced at the **fps** setting by retransmitting frames if the requested period is exceeded. This is sometimes needed in the case of an external trigger because PG cameras sometimes just ignore them. I have opened a support ticket on this, but PG has no solution yet.

TYPE = file

- **fps=float** Target frame rate in frames per second. If left undefined, frames will be read as quickly as possible.
- **roi={x_offset=+int, y_offset=+int, width=+int, height=+int}** Region of interest to extract from the camera or video stream (pixels).

TYPE = wcam

- **index=+int** User specified camera index. Useful in multi-camera imaging configurations.

Examples

```
# Serve to the 'wraw' stream from a webcam  
oat frameserve wcam wraw
```

```
# Stream to the 'graw' stream from a point-grey GIGE camera  
# using the gige_config tag from the config.toml file  
oat frameserve gige graw -c config.toml gige_config
```

```
# Serve to the 'fraw' stream from a previously recorded file  
# using the file_config tag from the config.toml file  
oat frameserve file fraw -f ./video.mpg -c config.toml file_config
```

Frame Filter

oat-framefilt - Receive frames from a frame source, filter, and publish to a second memory segment. Generally used to pre-process frames prior to object position detection. For instance, **framefilt** could be used to perform background subtraction or application of a mask to isolate a region of interest.

Signature

```
frame --> oat-framefilt --> frame
```

Usage

Usage: **framefilt** [INFO]

or: **framefilt** TYPE SOURCE SINK [CONFIGURATION]

Filter frames from SOURCE and published filtered frames to SINK.

TYPE

bsub: Background subtraction

mask: Binary mask

mog: Mixture of Gaussians background segmentation (Zivkovic, 2004)

undistort: Compensate for lens distortion using distortion model.

SOURCE:

User-supplied name of the memory segment to receive frames from (e.g. raw).

SINK:

User-supplied name of the memory segment to publish frames to (e.g. filt).

INFO:

--help Produce help message.

-v [--version] Print version information.

CONFIGURATION:

-c [--config] arg Configuration file/key pair.

-m [--invert-mask] If using TYPE=mask, invert the mask before applying

Configuration File Options

TYPE = bsub

- **background=string** Path to a background image to be subtracted from the SOURCE frames. This image must have the same dimensions as frames from SOURCE.

TYPE = mask

- **mask=string** Path to a binary image used to mask frames from SOURCE. SOURCE frame pixels with indices corresponding to non-zero value pixels in the mask image will be unaffected. Others will be set to zero. This image must have the same dimensions as frames from SOURCE.

TYPE = mog

- **learning_coeff=+float** Value, 0 to 1.0, specifying how quickly the statistical model of the background image should be updated. Default is 0, specifying no adaptation.
- **gpu_index=+int** Index of the GPU to use for performing background subtraction if Oat was compiled with CUDA support.

TYPE = undistort

- **camera-model=+int** Value, 0 or 1, specifying the camera model to use. 0 specifies a Pinhole camera model, 1 specifies fisheye. Generated by oat-calibrate.
- **camera-matrix= [+float,+float,+float,+float,+float,+float,float,+float,+float]**, Camera matrix for your imaging setup. Generated by oat-calibrate.
- **distortion-coeffs= [+float,+float,+float,+float,+float, ...]**, Five to eight element vector specifying lens distortion coefficients. Generated by oat-calibrate.
- **rotation=+double** Counter clockwise Degrees that undistorted image should be rotated. If not specified, defaults to 0.0.

Examples

```
# Receive frames from 'raw' stream
# Perform background subtraction using the first frame as the background
# Publish result to 'sub' stream
oat framefilt bsub raw sub

# Receive frames from 'raw' stream
# Apply a mask specified in a configuration file
# Publish result to 'roi' stream
oat framefilt mask raw roi -c config.toml mask-config
```

Frame Viewer

oat-view - Receive frames from named shared memory and display them on a monitor. Additionally, allow the user to take snapshots of the currently displayed frame by pressing **s** while the display window is in focus.

Signature

frame --> oat-view

Usage

Usage: view [INFO]
or: view SOURCE [CONFIGURATION]
Display frame SOURCE on a monitor.

SOURCE:

User-supplied name of the memory segment to receive frames from (e.g. raw).

INFO:

--help Produce help message.
-v [--version] Print version information.

CONFIGURATION:

-n [--filename] arg The base snapshot file name.
 The timestamp of the snapshot will be prepended to
 this name.If not provided, the SOURCE name will be
 used.

-f [--folder] arg The folder in which snapshots will be saved.

Example

```
# View frame stream named raw  
oat view raw
```



```
# View frame stream named raw and specify that snapshots should be saved  
# to the Desktop with base name 'snapshot'  
oat view raw -f ~/Desktop -n snapshot
```

Position Detector

`oat-posidet` - Receive frames from named shared memory and perform object position detection within a frame stream using one of several methods. Publish detected positions to a second segment of shared memory.

Signature

`frame --> oat-posidet --> position`

Usage

Usage: `posidet` [INFO]
or: `posidet` TYPE SOURCE SINK [CONFIGURATION]
Perform object position detection on frames from SOURCE.
Publish detected object positions to SINK.

TYPE
 `diff`: Difference detector (grey-scale, motion)
 `hsv` : HSV detector (color)

SOURCE:
 User-supplied name of the memory segment to receive frames from (e.g. `raw`).

SINK:
 User-supplied name of the memory segment to publish detected positions to (e.g. `pos`).

INFO:
 `--help` Produce help message.
 `-v [--version]` Print version information.

CONFIGURATION:
 `--tune` Use GUI to tune detection parameters at the cost of performance.
 `-c [--config] arg` Configuration file/key pair.

Configuration File Options

TYPE = `hsv`

- `tune=bool` Provide GUI sliders for tuning hsv parameters

- **erode**=+int Candidate object erosion kernel size (pixels)
- **dilate**=+int Candidate object dilation kernel size (pixels)
- **min_area**=+double Minimum object area (pixels2)
- **max_area**=+double Maximum object area (pixels2)
- **h_thresholds**=+{min=+int, max=+int} Hue pass band
- **s_thresholds**=+{min=+int, max=+int} Saturation pass band
- **v_thresholds**=+{min=+int, max=+int} Value pass band

TYPE = diff

- **tune**=bool Provide GUI sliders for tuning diff parameters
- **blur**=+int Blurring kernel size (normalized box filter; pixels)
- **diff_threshold**=+int Intensity difference threshold

Example

```
# Use color-based object detection on the 'raw' frame stream
# publish the result to the 'cpos' position stream
# Use detector settings supplied by the hsv_config key in config.toml
oat posidet hsv raw cpos -c config.toml hsv_config

# Use motion-based object detection on the 'raw' frame stream
# publish the result to the 'mpos' position stream
oat posidet diff raw mpos
```

Position Generator

`oat-posigen` - Generate positions for testing downstream components. Publish generated positions to shared memory.

Signature

`oat-posigen --> position`

Usage

Usage: `posigen` [INFO]
or: `posigen` TYPE SINK [CONFIGURATION]
Publish generated positions to SINK.

TYPE
 `rand2D`: Randomly accelerating 2D Position

SINK:
 User supplied position sink name (e.g. `pos`).

OPTIONAL ARGUMENTS:

INFO:
 `--help` Produce help message.
 `-v [--version]` Print version information.

CONFIGURATION:
 `-r [--sps] arg` Samples per second. Overridden by information in configuration file if provided.
 `-c [--config] arg` Configuration file/key pair.

Configuration File Options

TYPE = `rand2D`

- **`dt=+double`** Position update period
- **`room=[+double, +double, +double, +double]`** The ‘room’ in which generated positions reside specified as [x origin, y origin, width, height]. Arbitrary units. The room has periodic boundaries so when a position leaves one side it will enter the opposing one.

Example

```
# Publish randomly moving positions to the 'pos' position stream  
oat posigen rand2D pos
```


Position Filter

oat-posifilt - Receive positions from named shared memory, filter, and publish to a second memory segment. Can be used to, for example, remove discontinuities due to noise or discontinuities in position detection with a Kalman filter or annotate categorical position information based on user supplied region contours.

Signature

position --> oat-posifilt --> position

Usage

Usage: posifilt [INFO]
or: posifilt TYPE SOURCE SINK [CONFIGURATION]
Filter positions from SOURCE and published filtered positions to SINK.

TYPE

kalman: Kalman filter
homography: homography transform
region: position region label annotation

SOURCE:

User-supplied name of the memory segment to receive positions from (e.g. rpos).

SINK:

User-supplied name of the memory segment to publish positions to (e.g. rpos).

INFO:

--help Produce help message.
-v [--version] Print version information.

CONFIGURATION:

-c [--config] arg Configuration file/key pair.

Configuration File Options

TYPE = kalman

- **dt=+float** Sample period (seconds).
- **timeout=+float** Time to perform position estimation detection with lack of updated position measure (seconds).

- **sigma_accel**=+float Standard deviation of normally distributed, random accelerations used by the internal model of object motion (position units/s²; e.g. pixels/s²).
- **sigma_noise**=+float Standard deviation of randomly distributed position measurement noise (position units; e.g. pixels).
- **tune**=bool Use the GUI to tweak parameters.

TYPE = homography

- **homography** = [+float,+float,+float, +float,+float,+float,+float,+float,+float], Homography matrix for 2D position (1x9; world units/pixel). Generate using oat-calibrate.

TYPE = region

- **<regions>**=[+float, +float],[+float, +float],...,[+float, +float]] User-named region contours (pixels). Regions counters are specified as n-point matrices, [[x0, y0],[x1, y1],...,[xn, yn]], which define the vertices of a polygon. The name of the contour is used as the region label. For example, here is an octagonal region called CN and a tetragonal region called R0:

These regions could be named anything...

```
CN = [[336.00, 272.50],
      [290.00, 310.00],
      [289.00, 369.50],
      [332.67, 417.33],
      [389.33, 413.33],
      [430.00, 375.33],
      [433.33, 319.33],
      [395.00, 272.00]]
```

```
R0 = [[654.00, 380.00],
      [717.33, 386.67],
      [714.00, 316.67],
      [655.33, 319.33]]
```

Example

```
# Perform Kalman filtering on object position from the 'pos' position stream
# publish the result to the 'kpos' position stream
# Use detector settings supplied by the kalman_config key in config.toml
oat posifilt kalman pos kfilt -c config.toml kalman_config
```

Position Combiner

oat-posicom - Combine positions according to a specified operation.

Signature

```
position 0 --> |
position 1 --> |
      :         | oat-posicom --> position
position N --> |
```

Usage

Usage: posicom [INFO]
or: posicom TYPE SOURCES SINK [CONFIGURATION]
Combine positional information from two or more SOURCES.
Publish combined position to SINK.

TYPE

mean: Geometric mean of SOURCE positions

SOURCES:

User-supplied position source names (e.g. pos1 pos2).

SINK:

User-supplied position sink name (e.g. pos).

INFO:

--help	Produce help message.
-v [--version]	Print version information.

CONFIGURATION:

-c [--config] arg	Configuration file/key pair.
---------------------	------------------------------

Configuration File Options

TYPE = mean

- **heading_anchor=+int** Index of the SOURCE position to use as an anchor when calculating object heading. In this case the heading equals the mean directional vector between this anchor position and all other SOURCE positions. If unspecified, the heading is not calculated.

Example

```
# Generate the geometric mean of 'pos1' and 'pos2' streams
# Publish the result to the 'com' stream
oat posicom mean pos1 pos2 com
```

Frame Decorator

oat-decorate - Annotate frames with sample times, dates, and/or positional information.

Signature

```
      frame --> |
position 0 --> |
position 1 --> | oat-decorate --> frame
      :         |
position N --> |
```

Usage

Usage: decorate [INFO]

or: decorate SOURCE SINK [CONFIGURATION]

Decorate the frames from SOURCE, e.g. with object position markers and sample number. Publish decorated frames to SINK.

SOURCE:

User-supplied name of the memory segment from which frames are received (e.g. raw).

SINK:

User-supplied name of the memory segment to publish frames to (e.g. out).

OPTIONS:

INFO:

--help	Produce help message.
-v [--version]	Print version information.

CONFIGURATION:

-p [--position-sources] arg	The name of position server(s) used to draw object position markers.
-------------------------------	--

<code>-t [--timestamp]</code>	Write the current date and time on each frame.
<code>-s [--sample]</code>	Write the frame sample number on each frame.
<code>-S [--sample-code]</code>	Write the binary encoded sample on the corner of each frame.
<code>-R [--region]</code>	Write region information on each frame if there is a position stream that contains it.

Example

```
# Add textual sample number to each frame from the 'raw' stream
oat decorate raw -s

# Add position markers to each frame from the 'raw' stream to indicate
# objection positions for the 'pos1' and 'pos2' streams
oat decorate raw -p pos1 pos2
```

Recorder

`oat-record` - Save frame and position streams to file.

- **frame** streams are compressed and saved as individual video files (H.264 compression format AVI file).
- **position** streams are combined into a single JSON file. Position files have the following structure:

```
{oat-version: X.X},
{header: {timestamp: YYYY-MM-DD-hh-mm-ss},
        {sample_rate_hz: X.X},
        {sources: [ID_1, ID_2, ..., ID_N]} }
{positions: [ [ID_1: position, ID_2: position, ..., ID_N: position ],
              [ID_1: position, ID_2: position, ..., ID_N: position ],
              [ID_1: position, ID_2: position, ..., ID_N: position ] }
}
```

where each **position** object is defined as:

```
{
  samp: Int,           | Sample number
  unit: Int,           | Enum specifying length units (0=pixels, 1=meters)
  pos_ok: Bool,        | Boolean indicating if position is valid
  pos_xy: [Double, Double], | Position x,y values
  vel_ok: Bool,        | Boolean indicating if velocity is valid
  vel_xy: [Double, Double], | Velocity x,y values
  head_ok: Bool,       | Boolean indicating if heading is valid
  head_xy: [Double, Double], | Heading x,y values
  reg_ok: Bool,        | Boolean indicating if region tag is valid
  reg: String          | Region tag
}
```

Data fields are only populated if the values are valid. For instance, in the case that only object position is valid, and the object velocity, heading, and region information are not calculated, an example position data point would look like this:

```
{ samp: 501,
  unit: 0,
  pos_ok: True,
  pos_xy: [300.0, 100.0],
  vel_ok: False,
  head_ok: False,
  reg_ok: False }
```

All streams are saved with a single recorder have the same base file name and save location (see usage). Of course, multiple recorders can be used in parallel to (1) parallelize the computational load of video compression, which tends to be quite intense and (2) save to multiple locations simultaneously.

Signature

```
position 0 --> |
position 1 --> |
:              |
position N --> | oat-record
                |
    frame 0 --> |
    frame 1 --> |
    :           |
    frame N --> |
```

Usage

```
Usage: record [OPTIONS]
       or: record [CONFIGURATION]
```

OPTIONS:

```
--help           Produce help message.
-v [ --version ] Print version information.
```

CONFIGURATION:

```
-n [ --filename ] arg  The base file name to which to source name will
                        be appended
-f [ --folder ] arg    The path to the folder to which the video
                        stream and position information will be saved.
                        If specified, YYYY-MM-DD-hh-mm-ss_ will be
                        prepended to the filename.
-p [ --positionsources ] arg The name of the server(s) that supply object
                        position information. The server(s) must be of
                        type SMServer<Position>
-i [ --imagesources ] arg The name of the server(s) that supplies images
                        to save to video. The server must be of type
                        SMServer<SharedCVMatHeader>
```

Example

```
# Save positional stream 'pos' to current directory
oat record -p pos

# Save positional stream 'pos1' and 'pos2' to current directory
oat record -p pos1 pos2

# Save positional stream 'pos1' and 'pos2' to Desktop directory and
# prepend the timestamp to the file name
oat record -p pos1 pos2 -d -f ~/Desktop

# Save frame stream 'raw' to current directory
oat record -i raw

# Save frame stream 'raw' and positional stream 'pos' to Desktop
# directory and prepend the timestamp and the word 'test' to each filename
oat record -i raw -p pos -d -f ~/Desktop -n test
```


Position Socket

oat-posisock - Stream detected object positions to the network in either client or server configurations.

Signature

position --> oat-posisock

Usage

Usage: posisock [INFO]
or: posisock TYPE SOURCE [ENDPOINT]
Send positions from SOURCE to a remote endpoint.

TYPE:

std: Asynchronous position dump to stdout.
pub: Asynchronous position publisher over ZMQ socket.
Publishes positions without request to potentially many subscribers.
rep: Synchronous position replier over ZMQ socket.
Sends positions in response to requests from a single endpoint. Several transport/protocol options. The most useful are tcp and interprocess (ipc).
udp: Asynchronous, client-side, unicast user datagram protocol over a traditional BSD-style socket.

ENDPOINT:

Device to send positions to.

When TYPE is pos or rep, this is specified using a ZMQ-style endpoint: '<transport>://<host>:<port>'. For instance, 'tcp://*:5555' or 'ipc://*:5556' specify TCP and interprocess communication on ports 5555 or 5556, respectively
When TYPE is udp, this is specified as '<host> <port>'
For instance, '10.0.0.1 5555'.

INFO:

--help Produce help message.
-v [--version] Print version information.

Example

```
# Stream positions from the 'pos' stream to port 5555 at 18.72.0.3 in  
# client mode  
oat posisock pos -h 18.72.0.3 -p 5555
```

Buffer

oat-buffer - A first in, first out (FIFO) token buffer that can be use to decouple asynchronous portions of a data processing network. An example of this is the case when a precise external clock is used to govern image acquisition via a physical trigger line. In this case, 'hickups' in the data processing network following the camera should not cause the camera to skip frames. Of course, there is no free lunch: if the processing pipline cannot keep up with the external clock on average, then the buffer will eventually fill and overflow.

Signatures

```
position --> oat-buffer --> position
```

```
frame --> oat-buffer --> frame
```

Usage

Usage: buffer [INFO]

or: buffer TYPE SOURCE SINK [CONFIGURATION]

Place tokens from SOURCE into a FIFO. Publish tokens in FIFO to SINK.

TYPE

frame: Frame buffer

pos2D: 2D Position buffer

SOURCE:

User-supplied name of the memory segment to receive tokens from (e.g. input).

SINK:

User-supplied name of the memory segment to publish tokens to (e.g. output).

OPTIONS:

INFO:

--help Produce help message.

-v [--version] Print version information.

Example

```
# Acquire frames on a gige camera driven by an exnternal trigger  
oat frameserve gige raw -c config.toml gige-trig
```

```

# Buffer the frames separate asynchronous sections of the processing network
oat buffer frame raw buff

# Filter the buffered frames and save
oat framefilt mog buff filt
oat record -f ~/Desktop/ -p buff filt

```

In the above example, one must be careful to fully separate the network across the buffer boundary in order for it to provide any functionality. For instance, if we changed the record command to the following

```

oat record -f ~/Desktop/ -p raw filt

```

Then the buffer would do nothing since the raw token stream must be synchronous with the recorder, which bypasses the buffer. In this case, the buffer is just wasting CPU cycles. Here is a graphical representation of the first configuration where the `oat-buffer` is used properly. The synchronization boundary is shown using vertical lines.

```

      |
frameserve --> buffer --> framefilt --> record
      |   \                      /
      |   -----

```

In the second configuration, the connection from frameserve to record breaks the synchronization boundary.

```

      |
frameserve --> buffer --> framefilt --> record
      \   |                      /
      ----|-----

```

Calibrate

oat-calibrate - Interactive program used to generate calibration parameters for an imaging system that can be used to parameterize **oat-framefilt** and **oat-posifilt**. Detailed usage instructions are displayed upon program startup.

Signature

```
frame --> oat-calibrate
```

Usage

```
Usage: calibrate [INFO]
       or: calibrate SOURCE [CONFIGURATION]
Generate camera calibration and homography transform for a frame SOURCE.
```

TYPE

camera: Generate calibration parameters (camera matrix and distortion coefficients).
homography: Generate homography transform between pixels and world units.

SOURCE:

User-supplied name of the memory segment to receive frames from (e.g. raw).

OPTIONS:

INFO:

```
--help          Produce help message.
-v [ --version ] Print version information.
```

CONFIGURATION:

`-n [--calibration-name] arg` The key name for the calibration that will be inserted into the calibration file.

-f [--calibration-path] arg	The base calibration file location. If not is specified, defaults to './calibration.toml'. If a folder is specified, defaults to ' <folder> /calibration.toml'
-------------------------------	--

--homography-method arg Homography estimation method.

Values:

```
robust (default): RANSAC-based robust
estimation method (automatic outlier
rejection).
```

	regular: Best-fit using all data points. exact: Compute the homography that fits four points. Useful when frames contain know fiducial marks.
<code>--camera-model arg</code>	Model used for camera calibration. Values: pinhole (default): Pinhole camera model. fisheye: Fisheye camera model (ultra wide-angle lens with pronounced radial distortion.
<code>-h [--chessboard-height] arg</code>	The number of vertical black squares in the chessboard used for calibration.
<code>-w [--chessboard-width] arg</code>	The number of horizontal black squares in the chessboard used for calibration.
<code>-W [--square-width] arg</code>	The length/width of a single chessboard square in meters.
<code>-c [--config] arg</code>	Configuration file/key pair.

Kill

`oat-kill` - Issue SIGINT to all running Oat processes started by the calling user. A side effect of Oat's architecture is that components can become orphaned in certain circumstances: abnormal termination of attached sources or sinks, running pure sources in the background and forgetting about them, etc. This utility will gracefully interrupt all currently running oat components.

Usage

Usage: `kill`

Example

```
# Interrupt all currently running oat components  
oat kill
```

Clean

oat-clean - Programmer's utility for cleaning shared memory segments after following abnormal component termination. Not required unless a program terminates without cleaning up shared memory. If you are using this for things other than development, then please submit a bug report.

Usage

Usage: clean [INFO]

or: clean NAMES [CONFIGURATION]

Remove the named shared memory segments specified by NAMES.

INFO:

<code>--help</code>	Produce help message.
<code>-v [--version]</code>	Print version information.
<code>-q [--quiet]</code>	Quiet mode. Prevent output text.
<code>-l [--legacy]</code>	Legacy mode. Append <code>"_sh_mem"</code> to input NAMES before removing.

Example

```
# Remove raw and filt blocks from shared memory after abnormal termination of  
# some components that created them  
oat clean raw filt
```


Installation

First, ensure that you have installed all dependencies required for the components and build configuration you are interested in using. For more information on dependencies, see the dependencies section below. To compile and install Oat, starting in the top project directory, create a build directory, navigate to it, and run cmake on the top-level CMakeLists.txt like so:

```
mkdir release
cd release
cmake -DCMAKE_BUILD_TYPE=Release [CMAKE OPTIONS] ..
make
make install
```

If you just want to build a single component component, individual components can be built using `make [component-name]`, e.g. `make oat-view`. Available cmake options and their default values are:

```
-DUSE_FLYCAP=Off // Compile with support for Point Grey Cameras
-DBUILD_DOCS=Off // Generate Doxygen documentation
```

If you had to install Boost from source, you must let cmake know where it is installed via the following switch. Obviously, provide the correct path to the installation on your system.

```
-DBOOST_ROOT=/opt/boost_1_59_0
```

To complete installation, add the following to your `.bashrc` or equivalent. This makes Oat commands available within your user profile (once you start a new terminal):

```
# Make Oat commands available to user
eval "$(path/to/Oat/oat/bin/oat init -)"
```

Dependencies

License compatibility

Oat is licensed under the GPLv3.0. Its dependencies' are licenses are shown below:

- Flycapture SDK: NON-FREE specialized license (This is an optional package. If you compile without Flycapture support, you can get around this. Also, see the `GigE interface cleanup` entry in the TODO section for a potentially free alternative.

- OpenCV: BSD
- ZeroMQ: LGPLv3.0
- Boost: Boost software license
- cptoml: Some kind of Public Domain Dedication
- RapidJSON: BSD
- Catch: Boost software license

These licenses do not violate the terms of Oat's license. If you feel otherwise please submit an bug report.

Flycapture SDK

The FlyCapture SDK is used to communicate with Point Grey digital cameras. It is not required to compile any Oat components. However, the Flycapture SDK is required if a Point Grey camera is to be used with the `oat-frameserve` component to acquire images. If you simply want to process pre-recorded files or use a web cam, e.g. via

```
oat-frameserve file raw -f video.mpg
oat-frameserve wcam raw
```

then this library is *not* required.

To install the Point Grey SDK:

- Go to point-grey website
- Download the FlyCapture2 SDK (version $\geq 2.7.3$). Annoyingly, this requires you to create an account with Point Grey.
- Extract the archive and use the `install_flycapture.sh` script to install the SDK on your computer and run

```
tar xf flycapture.tar.gz
cd flycapture
sudo ./install_flycapture
```

Boost

The Boost libraries are required to compile all Oat components. You will need to install versions ≥ 1.56 . To install Boost, use APT or equivalent,

```
sudo apt-get install libboost-all-dev
```

If you are using an Ubuntu distribution older than Wily Werewolf, Boost will be too old and you will need to install from source via

```
# Install latest boost
wget http://sourceforge.net/projects/boost/files/latest/download?source=files -O tarboost
tar -xf tarboost
cd ./boost*
./bootstrap.sh
./b2 --with-program_options --with-system --with-thread
cd ..
sudo mv boost* /opt
```

OpenCV

opencv is required to compile the following oat components:

- oat-frameserve
- oat-framefilt
- oat-view
- oat-record
- oat-posidet
- oat-posifilt
- oat-decorate
- oat-positest

Note: OpenCV must be installed with ffmpeg support in order for offline analysis of pre-recorded videos to occur at arbitrary frame rates. If it is not, gstreamer will be used to serve from video files at the rate the files were recorded. No cmake flags are required to configure the build to use ffmpeg. OpenCV will be built with ffmpeg support if something like

```
-- FFMPEG:          YES
-- codec:           YES (ver 54.35.0)
-- format:          YES (ver 54.20.4)
-- util:            YES (ver 52.3.0)
-- swscale:         YES (ver 2.1.1)
```

appears in the cmake output text. The dependencies required to compile OpenCV with ffmpeg support, can be obtained as follows:

TODO

Note: To increase Oat's video visualization performance using `oat view`, you can build OpenCV with OpenGL and/or OpenCL support. Both will open up significant processing bandwidth to other Oat components and make for faster processing pipelines. To compile OpenCV with OpenGL and OpenCL support, first install dependencies:

```
sudo apt-get install libgtkglext1 libgtkglext1-dev
```

Then, add the `-DWITH_OPENGL=ON` and the `-DWITH_OPENCL=ON` flags to the cmake command below. OpenCV will be build with OpenGL and OpenCL support if `OpenGL support: YES` and `Use OpenCL: YES` appear in the cmake output text. If OpenCV is compiled with OpenCL and OpenGL support, the performance benefits will be automatic, no compiler options need to be set for Oat.

Note: If you have NVIDIA GPU that supports CUDA, you can build OpenCV with CUDA support to enable GPU accelerated video processing. To do this, will first need to install the CUDA toolkit.

- Be sure to **carefully** read the installation instructions since it is a multistep process. Here are some additional hints that worked for me:
- I have found that installing the toolkit via 'runfile' to be the most painless. To do this you will need to switch your system to text mode using `Ctrl + Alt + F1`, and killing the X-server via `sudo service lightdm stop` (or equivalent), and running the runfile with root privileges.
- I have had the most success on systems that do not use GNOME or other fancy desktop environments. The install on lubunut, which uses LXDE as its desktop environment, was especially smooth.
- Do **not** install the nvidia drivers along with the CUDA toolkit installation. I found that (using ubuntu 14.04) this causes all sorts of issues with X, cinnamon, etc, to the point where I could not even boot my computer into anything but text mode. Instead, install the NVIDIA drivers using either the package manager (`nvidia-current`) or even more preferably, using the `[device-drivers]`'(<http://askubuntu.com/a/476659>) program or equivalent.
- If you hare getting a `cv::exception` complaining that about `code=30(cudaErrorUnknown) "cudaGetDeviceCount(&device_count)"` or similar, run the affected command as root one time.

If OpenCV is compiled with CUDA suport, the CUDA-enabled portions of the Oat codebase will be enabled automatically. No compile flags are required.

Note: GUI functionality is enhanced in OpenCV is compiled with Qt support. You can build OpenCV with Qt by first installing the Qt SDK and these dependencies:

```
# Additional dependencies for integraged QT with OpenGL
sudo apt-get install libqt5opengl5 libqt5opengl5-dev
```

The you can compile OpenCV using QT support by adding `-DWITH_QT=ON` flag to the cmake command below. QT functionality will then be used by Oat automatically.

Finally, to compile and install OpenCV:

```

# Install OpenCV's dependencies
sudo apt-get install build-essential # Compiler
sudo apt-get install cmake git libgtk2.0-dev pkg-config libavcodec-dev libavformat-dev libsw
sudo apt-get install python-dev python-numpy libtbb2 libtbb-dev libjpeg-dev libpng-dev libt
sudo apt-get install # ffmpeg support [TODO]
sudo apt-get install # OpenGL support [TODO]
sudo ldconfig -v

# Get OpenCV
wget https://github.com/Itseez/opencv/archive/3.0.0-rc1.zip -O opencv.zip
unzip opencv.zip -d opencv

# Build OpenCV
cd opencv/opencv-3.0.0-rc1
mkdir release
cd release

# Run cmake to generate Makefile
# Add -DWITH_CUDA=ON for CUDA support and -DWITH_OPENGL for OpenGL support
cmake -DCMAKE_BUILD_TYPE=RELEASE -DCMAKE_INSTALL_PREFIX=/usr/local ..

# Build the project and install
make
sudo make install

```

ZeroMQ

ZeroMQ is required by the following Oat components:

- oat-record
- oat-posisock

TODO

RapidJSON, cpptoml, and Catch

These libraries are installed automatically by cmake during the build process.

RapidJSON is required by the following Oat components:

- oat-record
- oat-posisock

cpptoml is required by the following Oat components:

- `oat-frameserve`
- `oat-framefilt`
- `oat-posidet`
- `oat-posifilt`
- `oat-posicom`
- `oat-positest`

Catch is required to make and run tests using `make test`

Setting up a Point-grey PGE camera in Linux

`oat-frameserve` supports using Point Grey GIGE cameras to collect frames. I found the setup process to be straightforward and robust, but only after cobbling together the following notes.

Camera IP Address Configuration

First, assign your camera a static IP address. The easiest way to do this is to use a Windows machine to run the Point Grey ‘GigE Configurator’. If someone knows a way to do this without Windows, please tell me. An example IP Configuration might be:

- Camera IP: 192.168.0.1
- Subnet mask: 255.255.255.0
- Default gateway: 192.168.0.64

Point Grey GigE Host Adapter Card Configuration

Using network manager or something similar, you must configure the IPv4 configuration of the GigE host adapter card you are using to interface the camera with your computer.

- First, set the ipv4 method to **manual**.
- Next, you must configure the interface to (1) have the same network prefix and (2) be on the same subnet as the camera you setup in the previous section.
 - Assuming you used the camera IP configuration specified above, your host adapter card should be assigned the following private IPv4 configuration:
 - * POE gigabit card IP: 192.168.0.100
 - * Subnet mask: 255.255.255.0

* DNS server IP: 192.168.0.1

- Next, you must enable jumbo frames on the network interface. Assuming that the camera is using `eth2`, then entering

```
sudo ifconfig eth2 mtu 9000
```

into the terminal will enable 9000 MB frames for the `eth2` adapter. - Finally, to prevent image tearing, you should increase the amount of memory Linux uses for network receive buffers using the `sysctl` interface by typing

```
sudo sysctl -w net.core.rmem_max=1048576 net.core.rmem_default=1048576
```

into the terminal. *In order for these changes to persist after system reboots, the following lines must be added to the bottom of the `/etc/sysctl.conf` file:*

```
net.core.rmem_max=1048576
net.core.rmem_default=1048576
```

These settings can then be reloaded after reboot using

```
sudo sysctl -p
```

Multiple Cameras

- If you have two or more cameras/host adapter cards, they can be configured as above but *must exist on a separate subnets*. For instance, we could repeat the above configuration steps for a second camera/host adapter card using the following settings:

- Camera Configuration:
 - * Camera IP: 192.168.1.1
 - * Subnet mask: 255.255.255.0
 - * Default gateway: 192.168.1.64
- Host adapter configuration:
 - * POE gigabit card IP: 192.168.1.100
 - * Subnet mask: 255.255.255.0
 - * DNS server IP: 192.168.1.1

Example Camera Configuration

Below is an example network adapter and camera configuration for a two-camera imaging system provided by Point Grey. It consists of two Blackfly GigE cameras (Point Grey part number: BFLY-PGE-09S2C) and a single dual-port POE GigE adapter card (Point Grey part number: GIGE-PCIE2-2P02).

Camera 0

- Adapter physical connection (looking at back of computer)

```
RJ45  -----  
      |         |  
L [   [ ]     [x]   ] R
```

- Adapter Settings
 - Model: Intel 82574L Gigabit Network Connection
 - MAC: 00:B0:9D:DB:D9:63
 - MTU: 9000
 - DHCP: Disabled
 - IP: 192.168.0.100
 - Subnet mask: 255.255.255.0
- Camera Settings
 - Model: Blackfly BFLY-PGE-09S2C
 - Serial No.: 14395177
 - IP: 192.168.0.1 (Static)
 - Subnet mask: 255.255.255.0
 - Default GW: 0.0.0.0
 - Persistent IP: Yes

Camera 1

- Adapter physical connection (looking at back of computer)

```
RJ45  -----  
      |         |  
L [   [x]     [ ]   ] R
```

- Adapter Settings
 - Model: Intel 82574L Gigabit Network Connection
 - MAC: 00:B0:9D:DB:A7:29
 - MTU: 9000
 - DHCP: Disabled
 - IP: 192.168.1.100
 - Subnet mask: 255.255.255.0
- Camera Settings

- Model: Blackfly BFLY-PGE-09S2C
- Serial No.:
- IP: 192.168.1.1 (Static)
- Subnet mask: 255.255.255.0
- Default GW: 0.0.0.0
- Persistent IP: Yes

TODO

- [] ~~Mac, Windows builds?~~
- EDIT: Wait till somebody cares enough to do it.
- [] Unit and stress testing
 - Unit tests for `libshmndf`
 - * ~~Nominal data types, T~~
 - * Specializations for `Frames`
 - Stress tests for data processing chains
 - * I need to come up with a series of scripts that configure and run components in odd and intensive, but legal, ways to ensure sample sychronization is maintained, graceful exits, etc
- [] GigE interface cleanup
 - ~~The `PGGigCam` class is a big mess. It has has tons of code redundancy.~~
 - * EDIT: A lot of this is due to the PG API. I've cleaned up a bit, but more would be a waste of time.
 - ~~`oat-frameserve` `gige` can wait indefinitely if the cameras use an external trigger and that trigger source stops before the process is interrupted. Need a timed wait there.~~
 - * EDIT: Fixed in `a0c97e56bbe66227b03dd9253fdff33f0550465b`
 - Should I be using the generic GenICam API instead of PG's non-standard API? e.g. Aravis.
 - ~~Additionally, it needs to be optimized for performance. Are their unnecessary copies of images being made during conversion from PG Image to `cv::Mat`? Can I employ some move casts to help?~~
 - * EDIT: `shmndf` takes care of this.
 - There are a couple examples of GigE interfaces in OpenCV targeting other 3rd party APIs: `modules/videoio/src/cap_giganetix.cpp` and `opencv/modules/videoio/src/cap_pvapi.cpp`. I don't know that these are great pieces of code, but I should at least use them for inspiration.
 - ~~`oat-frameserve gige` lacks the ability to set FPS in free running (non-triggered mode)~~
 - ~~Configuration of the camera can get into impossible states if `oat` is used in combo with other programs that mess with the camera's registers. Configuration via `Oat` should start with a clean slate by setting the camera to a default register state.~~
 - ~~See `flycap` → advanced camera settings → restor default memory channel for how.~~
 - Binning should be specified in terms of mode number (see pg. 66 of blackfly tech ref) instead of bin size since not all bin sizes are allowed.

- [] Position type generalization
 - It might be a good idea to generalize the concept of a position to a multi-positional element
 - For things like the decorator, position combiner, and potentially detector, this could increase performance and decrease user script complexity if multiple targets common detection features needed to be tracked at once.
 - Down side is that it potentially increases code complexity and would require a significant refactor.
 - Additionally, position detection might no longer be stateless. E.g. think of the case when two detected objects cross paths. In order to ID the objects correctly in subsequent detections, the path of the objects would need to be taken into account (and there is not guarantee this result will be correct...). A potential work around is to have IDed ‘position groups’ with anonymous position members. This would get us back to stateless detection. However, it would make the concept of position combining hard to define (although that is even true now is just a design choice, really).
- [] Frame visualization is nice, but should no cause holdups in the processing pipeline. In this case frames should just be dropped. This might be as simple as making the actual visualization call in Viewer.cpp a call-back operating on a separate visualization thread that is triggered after `internal_frame_` is written to.
 - This may cause `internal_frame_` to be overwritten during visualization. Problem?
 - Visualization could somehow be ‘aborted’ if it has not completed and a new frame is available.
- [] Saving tuning parameters
 - Components that have a `--tune` option should also allow for the user to press a key and those tuning paramters to be injected into the current `config.toml` file so that they don’t have to write them down and manually edit the file later
- [] Colors
 - ~~Should visual ID information (e.g. color) be integrated into the `position` type:~~
 - * EDIT: No. What does position have to do with color? Sound? Touch? or ...
 - All this output color formatting might be very stupid because it screws up log files when stdout or stderr are piped to file.
- [] It would be good to warn people if the inputs to a multisource component have different sample rates because the output rate will be enforced by the slowest source.

- [x] Something is wrong with sample synchronization
 - When working with Jennie’s data, we found that position samples were being recorded multiple times - they had the same sample number and position info. Seems to be very intermittent, but points to a serious issue with sample synchronization. It seems likely this occurring in the recorder component due to its mulitthreaded implementation.
 - EDIT: With `libshmemdf`, sample numbers now travel with samples. I need to use this to implement `asserts` in each component that (1) check that samples increase monotonically (buffer overflows means that samples can be skipped, so uniary incrementation is not a given) and (2) that multisource components are dealing with sychornized sample numbers when pull-based sychornization strategy is enforced (no external clock driving acquisition, so no chance for buffer overrun).
- [] Configuration via file
 - There are a couple issues with configuration via TOML file as it stands
 1. Command line switches should take precedence over TOML file options. This is standard practice, but is not how Oat works currently.
 2. For (all?) most components, `configure` is pure abstract in the component’s base class. This doesn’t make too much sense because options are often common to many components. For instance, in `oat-posigen`, the sample period, and number of samples parameters are certainly relevant to any implementation of the position generator idea. Therefore, `configure` should be abstract with a base implmentation containing guaranteed-to-be-common parameters.